Machine Learning vs. Deep Learning
A Real-World Case Study Comparison

Machine learning (ML) and deep learning (DL) are both subfields of artificial intelligence. Still, they differ significantly in how they process data and the types of problems they are best suited to solve. This report examines one real-world example of traditional machine learning and one real-world example of deep learning, both drawn from the financial services industry, to explain why each approach fits its respective problem and why the opposite approach would be a poor fit in each case.

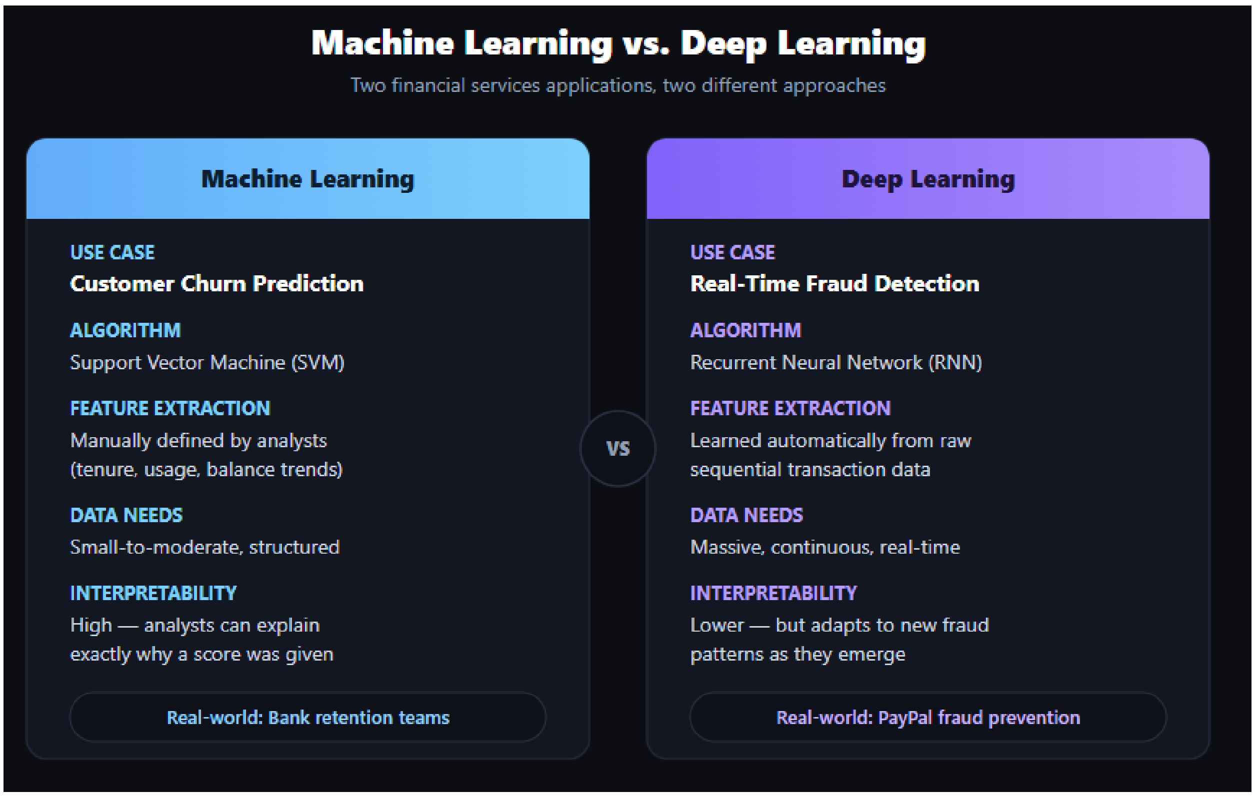

Example 1 - Machine Learning: Customer Churn Prediction

Real-World Application

Banks and financial services companies commonly use machine learning to predict customer churn, identifying customers who are likely to close their accounts or switch providers. A support vector machine (SVM) model can be trained on structured customer data such as account age, transaction frequency, product usage, and monthly balance trends to flag customers at high risk of leaving. Financial institutions use these predictions to trigger targeted retention offers, proactive outreach from relationship managers, or personalized incentives before a customer actually churns, which directly protects revenue and reduces the cost of acquiring new customers to replace lost ones.

Why Machine Learning Is Suitable here?

Customer churn prediction in banking is well-suited to traditional machine learning because the relevant features are structured, well understood, and can be defined by domain experts in advance, such as account tenure, number of products held, transaction frequency, and monthly charges. A support vector machine can efficiently learn from this kind of labeled, tabular historical data (customers who churned versus those who stayed) without requiring the massive data volumes or computing power that deep learning demands. This aligns with a key distinction between the two approaches: in machine learning, humans typically define the relevant features up front, while deep learning would instead need to learn those features automatically from much larger and less structured data. Because the feature set here is limited and interpretable, an SVM model also allows analysts to understand which factors are driving churn risk for a given customer, which is valuable for designing targeted retention strategies.

Why Deep Learning Is Not Suitable Here?

Deep learning would be a poor fit for this specific problem. Deep neural networks are designed to automatically extract complex, hierarchical patterns from large, unstructured datasets, which requires significantly more training data and computational power than a well-defined, structured task like churn prediction actually needs. Even though a deep learning model could theoretically be applied to this structured customer data, doing so would add unnecessary complexity and training cost for a problem that a simpler algorithm already solves efficiently and transparently. Since churn prediction does not require learning the kind of subtle, high-dimensional patterns found in raw sequential or unstructured data, the added complexity of deep learning would not meaningfully improve results and would instead make the model slower to train and harder for analysts to interpret when explaining retention decisions to stakeholders.

Example 2 - Deep Learning: Real-Time Payment Fraud Detection

Real-World Application

PayPal is a well-documented real-world application of deep learning for fraud detection. PayPal's fraud prevention platform uses recurrent neural networks (RNNs) to analyze sequences of a customer's transactions over time, learning to recognize the normal rhythm of a user's spending behavior so it can flag sudden deviations, such as unusual transaction velocity, atypical amounts, or suspicious geographic patterns, as they happen. This deep learning system reportedly analyzes hundreds of signals per transaction and has helped PayPal block billions of dollars in fraudulent activity annually, processing millions of transactions per hour in real time.

Why Deep Learning Is Suitable?

Real-time payment fraud detection is well-suited to deep learning because the problem involves analyzing continuous, sequential transaction data at a massive scale, where fraud patterns are subtle, constantly evolving, and not easily captured by a fixed set of hand-picked rules. A recurrent neural network is specifically designed to process sequences, much like it would process a sequence of words in machine translation, allowing it to maintain a kind of internal memory of a customer's past transaction behavior and automatically learn what a 'normal' spending pattern looks like for that individual. This automatic feature learning is critical because fraud tactics change constantly, and a model that can adapt directly from raw sequential transaction data will catch new fraud patterns that a fixed, manually engineered rule set would miss. Given the enormous volume of transactions processed every hour and the heavy computational investment financial institutions make in fraud prevention, deep learning's ability to model complex temporal patterns makes it a more effective approach for this problem.

Why Traditional Machine Learning Is Not Suitable Here?

Traditional machine learning would be a weaker fit for large-scale, real-time fraud detection. ML algorithms like logistic regression or decision trees rely on a human manually defining a limited set of relevant features ahead of time, but fraud patterns evolve too quickly and unpredictably for a fixed feature set to stay effective for long. A simpler ML model also struggles to capture the sequential, time-dependent nature of transaction behavior the way an RNN can, since it does not inherently account for the order and timing of past transactions when evaluating a new one. As fraud tactics shift, an ML model built on static, hand-engineered features would require constant manual retooling, whereas a deep learning model can continue learning evolving patterns directly from the raw sequential data, making it far better suited to a fast-moving, adversarial problem like payment fraud.

Conclusion

These two examples illustrate the core distinction between machine learning and deep learning within a single industry: financial services. As the lesson's own analogy puts it, machine learning is like giving a student a checklist of things to look for to make a decision, while deep learning is like teaching the student how to read and letting them figure things out based on experience. Customer churn prediction is well-suited to traditional ML because its relevant features are limited, well understood, and can be manually defined from structured account data, while real-time fraud detection requires deep learning because the volume, speed, and constantly evolving nature of sequential transaction data make manual feature engineering impractical. Choosing the right approach ultimately depends on matching the complexity of the data and the problem to the strengths of each technique, rather than assuming one method is universally superior to the other.

References

PayPal. (n.d.). *Machine learning fraud detection technologies*. PayPal US. https://www.paypal.com/us/brc/article/payment-fraud-detection-machine-learning

Jurgovsky, J., Granitzer, M., Ziegler, K., Calabretto, S., Portier, P.-E., He-Guelton, L., & Caelen, O. (2018). Sequence classification for credit-card fraud detection. *Expert Systems with Applications*, *100*, 234–245. https://doi.org/10.1016/j.eswa.2018.01.037